

Week 3 Lab Session

Nalin Bhatt, Ethan Kozlowski, Max Zhu

Adapted from Adam Wu, Wonje Yun (Spring 2024)

April 11, 2025

University of Chicago

1. Review
2. Assignment Submission Example
3. Review of Numba/MPI
4. GPU Jobs
5. Assignment Walk-through

- Github setup
- Midway code editing
- Set up the necessary libraries at the base environment

Assignment Submission Example

- `git clone` assignment repository to Midway3
- Create folders so that assignments are clear
- Create separate python files if necessary (do not have to attempt to answer all in one file)
- Write down your code, run the code and get results
- Write down the README as described in Canvas < Assignment 1 and link

Review of Numba/MPI/sbatch

- In pyCC, function signs are crucial. Need to figure out your input and output type.
- What happens when `mpirun` a job and between `comm.Scatter` and `comm.Gather`?
 - Unless `rank` is specified, every node receives the input and returns output
 - To have only one output, put the `print` or `return` under

```
1 if rank == 0:  
2     # YOUR CODE HERE
```

- `.sbatch` file contains the configuration of HPC hardware the user is requesting, and the command that it is requested to run.

For more SLURM commands see also: [Slurm documentation](#)

MPI for Python - 1

- `MPI.COMM_WORLD`: set up communication for mpi
- `comm.Get_size()`: return the number of processes
- `comm.Get_rank()`: return the rank of this process
- `comm.Bcast()`: broadcast a message from one process to all other processes in a group
- `comm.Scatter()`: scatter data from one process to all other processes in a group
- `comm.Gather()`: gather together values from a group of processes

For more explanation: [mpi4py.MPI documentation](#) and [bcast/scatter/gather example](#)

- `comm.Bcast(buf, root=0)`: broadcast the data(`buf`) at rank 0 to all workers
- `comm.Scatter(sendbuf, recvbuf, root=0)`: scatter data(`sendbuf`) at rank 0 to `recvbuf` at each workers.
- `comm.Gather(sendbuf, recvbuf, root=0)`: gather the data(`sendbuf`) in each workers to the `recvbuf` at rank 0.
- If you want to do some global-like operation before scattering or after gathering, the operation should go under `"if rank == 0:"` statement.

Using GPU with OpenCL

- `pyopencl` gives a pythonic access to OpenCL API.
- As in `mpi4py`, we need codes within the python function to make it run as an `opencl` job.
- `cl.create_some_context()`: create a `opencl` context 'somehow' (a device is chosen in an implementation-defined manner).
 - The `Context` is a construct that manages the set of devices, memory and command-queues.
- `cl.CommandQueue()`: create a command queue using the specified context

```
1 import pyopencl as cl
2 ctx = cl.create_some_context()
3 queue = cl.CommandQueue(ctx)
```

- `pyopencl.array.Array`: The `numpy.ndarray` work-alike, so that it functions like an numpy array but operates on OpenCL.
[pyopencl.array.Array](#)
- `pyopencl.elementwise.ElementwiseKernel`: Performs element-wise computation on the `pyopencl.array.Array` object.
 - context: the OpenCL context
 - arguments: the variable argument of the operation (need to assign data types, for arrays add * before the variable)
 - operation: the element-wise operation to perform
 - name: the function name of the compiled kernel
 - Ref: [ElementwiseKernel](#)

ElementwiseKernel Code Example

```
1 import pyopencl.array as cl_array
2 from pyopencl.elementwise import ElementwiseKernel
3
4 ctx = cl.create_some_context()
5 queue = cl.CommandQueue(ctx)
6
7 a_g = cl.array.to_device(queue, a_np)
8 b_g = cl.array.to_device(queue, b_np)
9
10 lin_comb = ElementwiseKernel(ctx,
11     "float k1, float *a_g, float k2, float *b_g, float *res_g",
12     "res_g[i] = k1 * a_g[i] + k2 * b_g[i]",
13     "lin_comb")
14
15 res_g = cl.array.empty_like(a_g)
16 lin_comb(2, a_g, 3, b_g, res_g)
```

- `pyopencl.reduction.ReductionKernel`: Performs reduction operation (eg. count, sum, etc.) on the `pyopencl.array.Array` object.
 - `context`: the OpenCL context
 - `dtypeout`: the dtype of the output
 - `neutral`: initial value
 - `map_expr`: the entry operation of the vector
 - `reduce_expr`: the operation to perform on the result of `map_expr`
 - `arguments`: the input (needs at least one argument to be vector)
 - Ref: [ReductionKernel](#)

ReductionKernel Code Example

```
1 import pyopencl.array as cl_array
2 from pyopencl.reduction import ReductionKernel
3
4 ctx = cl.create_some_context()
5 queue = cl.CommandQueue(ctx)
6
7 a = cl_array.arange(queue, 400, dtype=numpy.float32)
8 b = cl_array.arange(queue, 400, dtype=numpy.float32)
9
10 krnl = ReductionKernel(ctx, numpy.float32, neutral="0",
11                        reduce_expr="a+b", map_expr="x[i]*y[i]",
12                        arguments="__global float *x, __global float *y")
13
14 my_dot_prod = krnl(a, b).get()
```

GPU sbatch Commands

For using GPUs, different sbatch commands are needed.

- **-time**: set a limit on the total run time of the job allocation
- **-nodes**: number of GPU nodes to use
- **-partition**: set gpu for Midway3
- **-gres**: the number of gpus to request (gpu:1 for 1 GPU usage)
- **-ntasks-per-node**: number of cores to drive each gpu
- **-mem-per-cpu**: minimum memory required per usable allocated CPU

For more, see also: [RCC user guide for SLURM](#) and [Slurm documentation](#)

GPU sbatch

```
1 #!/bin/bash
2
3 #SBATCH --job-name=gpu_mpi
4 #SBATCH --output=gpu_mpi.out
5 #SBATCH --error=gpu_mpi.err
6 #SBATCH --nodes=1 # using 1 GPU node
7 #SBATCH --ntasks-per-node=1 # 1 CPU node to drive the GPU
8 #SBATCH --partition=gpu # using the GPU on Midway3
9 #SBATCH --gres=gpu:1 # requesting only one GPU
10 #SBATCH --account=macs30123
11
12 module load cuda python
13
14 mpirun python3 ./<YOUR_FILE_NAME>.py
15
```

Question 1,2 Review

- Question 1 (a) explicitly asks you to use pyCC, not jit.
- To make the code running, it would be helpful to have the below code.

```
1 if __name__ == "__main__":  
2     # YOUR CODE HERE
```

- Question 1 (b) asks you to compare the time speed comparison of up to using 20 cores, not more and not less.
- You should call the maximum number of nodes you are going to use in the sbatch.

Question 3 Objective

- **Objective:** compare the computation time and result between CPU and GPU
- Question (a): compare the computation time and result between CPU and GPU.
- Question (c): compare the computation time when the image is increased 10 times and 20 times.

Question 3 Cautions

- The main framework code and expected image of Q3 is given in the assignment description. You only need to make it into a code that can be run in OpenCL.
- The folder for accessing the input data are located in `/project2/mac30123/landsat8/`.
- **Only read the files in this directory, do not write on it.**
- **Restrict your CPU memory usage by 30G:** use `#SBATCH --mem-per-cpu=30G`
- **They are all in the assignment pdf, so please read carefully for the instructions before you run the code**

References i



Lisandro Dalcin. "MPI for Python". Accessed April 5, 2024.

<https://mpi4py.readthedocs.io/en/stable/index.html>



Khronos OpenCL Working Group. "The OpenCL Specification." Accessed April 5, 2024.

https://registry.khronos.org/OpenCL/specs/3.0-unified/html/OpenCL_API.html



PyOpenCL. "pyopencl 2024.1 documentation." Accessed April 5, 2024. <https://document.tician.de/pyopencl/>



Research Computing Center. "RCC User Guide." Accessed March 20, 2024. <https://rcc-uchicago.github.io/user-guide/>



SCHEDMD. "Slurm Workload Manager - sbatch." Accessed April 5, 2024. <https://slurm.schedmd.com/sbatch.html/>